

BEACON

A Benchmark for Efficient and Accurate Counting of Subgraphs

Xiangju Zhu¹ **Mohammad Matin Najafi² (presenting)** Chrysanthi Kosyfaki³

Xiaodong Li⁴ Reynold Cheng¹ Laks V. S. Lakshmanan⁵

ICDE 2026 | Experiment, Analysis & Benchmark Track

¹University of Hong Kong ²Huawei HK Research Center ³HKUST

⁴Xiamen University ⁵University of British Columbia


What is subgraph counting?

Subgraph counting in one picture

Given a host graph G and a query pattern q , count every subset of G 's vertices whose induced subgraph is isomorphic to q .


Host graph G

6 nodes · 8 edges



Query pattern q


a triangle (3-clique)



$$\text{count}(G, q) = |\{ S \subseteq V(G) : G[S] \cong q \}|$$

Result $N_q(G) = 4$

four triangles: $\{1,2,3\}$, $\{1,2,4\}$, $\{1,3,4\}$, $\{2,3,4\}$



■ $\{1,2,3\}$ ■ $\{1,2,4\}$ ■ $\{1,3,4\}$ ■ $\{2,3,4\}$

Why is it important?

Where does subgraph counting show up in the real world?

A handful of structural motifs explain a huge fraction of valuable signal across domains.



Social networks [17]

Facebook, LinkedIn, Twitter



Triangle — tight-knit friend group, recommendation signal

Triangle counts power Facebook's **"People You May Know"**: more shared triangles around a candidate $\Rightarrow$ closer community $\Rightarrow$ more clicks.



Biological networks [18–20]

Protein interaction, gene regulation



Feed-forward loop — filters noise, robust regulation

Counts of FFL motifs ($\text{FGF} \rightarrow \text{AKT} \rightarrow \text{mTOR}$) reveal which gene-regulatory circuits filter cellular noise — a predictor of disease robustness.



Financial networks [21]

Bank transfers, crypto, AML



Directed 3-cycle — money round-trip / wash trades

Counts of directed cycles of length 3–8 are the canonical AML feature for spotting layering in correspondent banking and crypto graphs (FATF 2023 typology report).

Why count and not detect?

Detection answers "does at least one exist?". **Counting** answers "how many?" — the magnitude is what differentiates noise, normal, and anomalous behaviour at scale.

Two camps — algorithmic vs ML

Two camps, complementary strengths — never compared head-to-head

Each pipeline hides the other's failure modes; BEACON puts them on the same scale.



Algorithmic methods

ESCAPE · EVOKE · SCOPE · MOTIVO



Pattern q + Graph G



Hand-crafted formulas /
decomposition



Exact count

- + Exact, deterministic; no training needed
- + Sub-second on small patterns at large scale
- One formula per pattern family; brittle to new structures
- Orbit count blows up at $k \geq 6$



ML-based methods

GIN · ID-GNN · GNN-AK⁺ · ESC-GNN · P-GNN · DeSCo · PPGN



Graph G + pattern q



GNN: ego-net / orbit / 2nd-
order



Approximate count

- + Single model, many queries
- + Amortises across many graphs / patterns
- + Naturally consumes node/edge labels
- Distribution shift $\Rightarrow$ accuracy collapses
- Pre-processing (ego-nets, distance enc.) often costs more than full exact counting

Where AL wins: small patterns ($k \leq 5$), homogeneous graphs, exactness required.

applies to small patterns; flips at $k \geq 6$

Where ML wins: large patterns ($k \geq 6$), labelled graphs, repeated queries.

Before BEACON: each paper picked its own datasets, output type, and metric — making honest comparison impossible.

Why a benchmark? Five blockers nobody had cleared

Five concrete blockers:

1. **No shared datasets.**
2. **Missing ground truth** — exact counts are expensive.
3. **Mismatched semantics** — induced vs. non-induced, local vs. global.
4. **Fragile environments** — GPU memory, deps, hyperparams differ.
5. **Weak baselines** — ML papers often compared only to outdated AL.


The pivot

We set out to build a new ML method.
We ended up building the **infrastructure** that lets anyone tell whether *any* method is genuinely better.

What the community gets

 Verified data  Docker images
 Controlled tests  Public
leaderboard

BEACON: 3 modules + 4 design principles



Four design principles (paper Sec. IV.C):



Unified & method-agnostic

Fair head-to-head between AL exact / approximate and ML.



Ground-truth first

Curated Oracle (26,435 graphs) with verified counts up to $k=5$.



Reproducible by construction


Docker pinned-deps + fixed CPU set; no “works on my machine”.



Controlled diversity & stress

BEACON-Sampler: explicit constraints on size, density, degree.

Oracle Dataset: structural diversity by design



Density slope ($|E|/|V|$) differs $5\times$ across domains.

- **26,435 graphs** — bioinformatics, vision, social, molecular.
- **All 29 patterns of size ≤ 5** counted exactly: induced & non-induced, local & global.
- **ID-constrained ground truth** (DeSCo-style ego-net counts) shipped to seed future work.
- **120+ features** per graph (diameter, density, clustering, max local count, ...).

BEACON-Sampler (PyPI: `rwdq`)

```
{nodes: [100,500], avg_deg: [1.5,5],  
local_triangle_max: [10000,20000]}
```

Who benefits — four researcher personas

Mirroring the paper's Fig. 3(b): four researcher personas, four lifecycle stages.

Bob explore

"How can I research deeply to find weaknesses and advantages of existing works?"

1. Pull baseline Dockers
2. Tweak runtime params; ablations
3. Try different sets / scenarios

Alex design

"How can I make my new SC algorithm easy to develop and convenient for users?"

1. Encapsulate dev env
2. Mount/install code in container
3. Publish container + usage

Candy test

"I have a preliminary algorithm — how do I test scalability, robustness, accuracy, efficiency?"

1. Sample with BEACON-Sampler
2. Test across aspects
3. Find weaknesses; iterate


David ship

"I have my final algorithm — how do I get a fair comparison with other methods?"

1. Pull competitor Dockers
2. Run all algs on shared sets
3. Report pros / cons

→ lifecycle: explore → design → prototype → ship →

Benchmark protocol + how we measure



29 patterns of size 3–5, plus star patterns up to 8 nodes.

Sets 1–4	Many small graphs (≤ 500 nodes), increasing density; Set 2 drops degree cap $\rightarrow$ power-law tails
Sets 5–8	Single large graph, increasing total size
Scenarios	Zero-shot / fine-tune / full retrain

How we measure (paper Defs. 9–10)

Q-error (multiplicative):

$$Q = \max\left(\frac{\hat{C}+1}{C+1}, \frac{C+1}{\hat{C}+1}\right)$$

1.0 = perfect; 10 = off by $10\times$. Insensitive to scale; great for rare patterns.

MAE (additive): $|\hat{C} - C|$

Critical for large counts (up to 10^{10}).

Always report both — small MAE + large Q-err means you fail on *rare* patterns.

Zero-shot is the realistic scenario

For a static graph you count once. Needing fine-tuning defeats the purpose of using ML.

Scope of this work — and why

✓ In scope (axis values)

Axis	Choice
Direction	Undirected
Edges	Induced
Input	Network-centric
Labels	Untyped (homogeneous)

What we test $k = 3..5$ (full) + k -stars to 8; **11 methods** (4 AL + 7 ML).

🚧 On the leaderboard roadmap

- Typed / labeled graphs (e.g. LSS [Zhao, VLDB'23])
- Directed graphs
- Dynamic / streaming
- Distributed counting

Why this scope?

- **Most valuable in practice** — Ribeiro et al. ACM CSUR'22 [8] surveys $k = 3..5$ as the practically-relevant range.
- **Common ground + tractable** — all 11 baselines support $k \leq 5$; orbits explode $58 \rightarrow 407 \rightarrow 4306$ at $k = 5, 6, 7$ (Sec. V.B).
- **Fair to AL** — labeled methods “rely on node labels as features and degenerate on unlabeled graphs” (Sec. II.e).

Experimental findings

Finding 1 🕒 Algorithmic methods dominate on small-pattern speed





Fig. 7. Runtime time (red “F” indicates failure) for Set 1-4.

- ML pre-processing (ego-nets, distance encodings) often exceeds full exact-method runtime.
- “F” = OOM — ESC-GNN, I²-GNN, PPGN fail on denser sets.
- GIN is the ML outlier: fastest ML at 4 s on Set 8.
- 🕒 For $k \leq 5$ ML is the wrong tool — use optimised AL indices.


Set 8 (single large graph):

SCOPE (exact)	114.5 s
MOTIVO	38.6 s
GIN	4.1 s
DeSCo	1607.0 s
Others	OOM


Finding 2 ☉ Accuracy collapses without fine-tuning



(a) Zero-Shot



(b) Fine-tuning



(c) Full Retraining

Fig. 6. Accuracy under different training scenarios on Set 1.

(a) Zero-shot

Most models lose to MOTIVO;
ID-GNN: Q-err < 3 on 5-cycles,
> 10^3 on 5-paths.

Fails on distribution shift.

(b) Fine-tuning

A few epochs $\rightarrow$ Q-err < 10^1 .

Needs *some* ground truth.

Training data must exist.


(c) Full retrain

Best accuracy but high overhead;
PPGN can *degrade* vs. fine-tune.


Defeats the point of amortising.

💡 ML needs substantial ground-truth to generalise — partially defeating its purpose for fast estimations.

Finding 3 Where ML systematically fails



(a) Graph Accuracy



(b) Node Accuracy

Fig. 8. Effect of graph size, density, node degree and clustering coefficient on accuracy on target 11.

Fine-tuned Q-err, target 13 ($\triangle + 2$ edges):

Algorithms	Set 1	Set 2	Set 3	Set 4
MOTIVO	108.88	469.86	2.43	1.13
DeSCo	1.93	1.39	1.13	1.21
GIN	8.33	24.63	7.69	2.23
GNN-AK⁺	4307.59	7612.56	173.02	1438.93
ID-GNN	4.45	26.48	1.12	1.76
ESC-GNN	29220.07	N/A	N/A	N/A
I²-GNN	5.01	N/A	N/A	N/A
PPGN	1.02	4.01	1.39	1.64

- High-degree, low-clustering graphs = worst case.
- Set 2: GNN-AK⁺ Q-err hits **7612**, ESC-GNN **29220**.
- DeSCo robust via ID-constrained ego-nets.
- PPGN (≥ 3 -WL) accurate but $O(n^3)$.
- 💡 Density heavily impacts ML; AL handles it gracefully.

Finding 4 ML takes the lead at large patterns

TABLE VI
STAR PATTERN ACCURACY WITH MAE AND Q-ERROR FOR STAR-SHAPED PATTERNS.

Alg	3-star		4-star		5-star		6-star		7-star		8-star	
	MAE	Q-error	MAE	Q-error	MAE	Q-error	MAE	Q-error	MAE	Q-error	MAE	Q-error
MOTIVO	1.1e+02	1.23	4.0e+03	1.72	1.1e+05	4.97	4.5e+07	2.80	7.6e+08	6.15	1.1e+10	87.0
DeSCo	1.7e+02	1.07	2.9e+03	1.07	4.4e+04	1.25	1.9e+07	1.35	5.0e+08	6.33	4.6e+09	10.4
GIN	2.8e+02	1.10	4.3e+03	1.13	3.8e+05	2.66	8.9e+06	38.0	1.3e+08	1.8e+03	1.5e+09	2.5e+04
GNN-AK ⁺	4.7e+02	1.21	2.0e+02	1.01	1.8e+04	1.48	5.8e+05	1.67	1.3e+07	1.12	1.4e+08	1.54
ID-GNN	3.4e+01	1.01	2.5e+03	1.10	7.0e+04	1.42	3.1e+06	1.92	5.3e+07	2.56	7.2e+08	9.00
ESC-GNN	2.2e+03	10.0	2.5e+04	44.4	3.1e+05	8.0e+02	8.3e+06	7.8e+03	1.2e+08	7.8e+03	1.5e+09	3.1e+04
I ² -GNN	5.5e-01	1.00	1.1e+01	1.01	1.8e+03	2.11	7.9e+04	1.01	2.9e+06	1.54	3.0e+07	1.36
PPGN	5.2e+02	1.65	5.1e+03	1.20	6.5e+03	1.22	4.9e+05	1.68	2.9e+06	1.22	5.9e+08	1.62

As star-pattern size grows from 3 to 8 nodes:

- **Exact AL infeasible:** orbit explodes $58 \rightarrow 407 \rightarrow 4306$.
- **I²-GNN** Q-error ≈ 1 ; $\text{GNN-AK}^+ < 2$.
- **MOTIVO:** Q-error $1.2 \rightarrow 87$ (3 $\rightarrow$ 8-star).
- **GIN** (1-WL): Q-error $\approx 25,000$ at 8-star.

The crossover point

$k \leq 5$:  dominates | $k \geq 6$:  often only option
BEACON locates this crossover.

Synthesis — who wins, where, and why

	Small graph	Large graph
Small pattern ($k \leq 5$)	 AL exact wins on <i>accuracy + speed</i> ESCAPE/EVOKE/SCOPE: < 1 s; ML pre-proc alone exceeds it	 AL exact / approximate wins Only GIN keeps up among ML (4 s on Set 8)
Large pattern ($k = 6 - 8$)	 ML wins; AL infeasible I^2 -GNN / GNN-AK ⁺ hold Q-err < 2	 Only ML scales PPGN OOMs; I^2 -GNN keeps Q-err ≈ 1

💡 The crossover lives at $k = 6$. **BEACON** is the first benchmark that actually *locates* it.

Next steps unlocked by BEACON

Hybrid AL + ML

Decompose with AL, learn the residual; “local counts of small patterns as features for larger patterns” (Sec. VI).

Helps Candy / David.

Typed / labeled track

Reactivate LSS [Zhao'23], Schwabe-Acosta [PACMMOD'24] on a labeled Oracle.

Helps Bob / David.

Directed graphs

Definition 1 currently assumes undirected; new Oracle counts needed.

Helps Alex.

Dynamic / streaming

Wackersreuther et al. MLG'10; no SOTA in benchmark yet.

Helps Candy.

Distributed counting

Zhang et al. VLDB'20 [39] is uncovered today.

Helps Alex.

Public live leaderboard

Ongoing deliverable; new methods plug in via Docker.

Helps everyone.

BEACON is open

Use it. Break it. Improve it.

 github.com/zxj0302/MLSC

 `pip install rwdq` — BEACON-Sampler on PyPI

 Docker Hub: pre-built images for all 11 baselines

Thank you.

 Leaderboard (**coming soon**): AL & ML on identical inputs

 matin@connect.hku.hk — questions, new datasets, contributions

Questions?